GraphSnapShot: An Efficient Framework for Graph Local Structure Snapshot Storage and Retrieval

November 2023

1 Introduction

In our recent research, we have developed a framework called GraphSnapShot, which has been proven an useful tool for graph learning acceleration. The core idea of GraphSnapShot is to capture and update the state of local graph structures dynamically, just like taking snapshots of graphs.

GraphSnapShot is designed to efficiently capture, store and update the dynamic snapshots of graph data, enabling us to track patterns in the structure of graph networks. This technique is useful for most graph learning tasks that relies on topology analysis or networks are constantly evolving, such as social media analysis, biological networks, or any system where the relationships between entities change over time.

The key components of GraphSnapShot is the GraphSDSampler. GraphSDSampler can efficiently capture, update, retrieve and store graph snapshots of topology while doing computation at the same time, which makes graph learning computation significantly faster.

In experiments, GraphSnapShot shows efficiency. It can promote computation speed by 20+ times compared to GraphSAGE in local topology storage and retrival (in 2-hop, 3-hop and 4-hop domain), with little loss of accuracy. It shows that the GraphSnapShot has potential to be a powerful tool for large graph acceleration.

2 Background

Local structure analysis in graph learning applications are fundamental for understanding and analyzing complex networks in various domains. This concept pertains to the idea of capturing the localized patterns and relationships within a large graph structure. Efficient graph structure learning are crucial in various applications, from social network analysis to bioinformatics, where useful hidden pattern are discovered by graph structure learning.

The GraphSnapShot framework we proposed provide insight to study graph learning by using dynamically updated local graph snapshots. GraphSnapShot is particularly valuable in large-scale networks where global analysis can be computationally intensive and less informative for certain types of inquiries. By focusing on local structures, researchers can detect community patterns, identify influential nodes, and understand local clustering behaviors in a more efficient way.

Applications of local structure learning are diverse and impact several fields:

- **Social Network Analysis:** Understanding individual or entity interactions within a network to reveal social dynamics and community formation.

- **Bioinformatics:** In biological networks, local structures assist in identifying functional modules or predicting protein interactions.

- **Recommendation Systems:** Improving recommendation accuracy by focusing on immediate user-item interaction patterns.

- **Network Security:** Analyzing local patterns for detecting anomalies or potential security breaches.

3 Motivation

Traditional sampling methods in graph analysis, particularly for multi-hop domains, struggle with efficiency. Node-wise sampling methods, including well-known ones like GraphSAGE [HYL18], are not optimized for multi-hop scenarios due to the computational burden that grows exponentially with each additional hop—a phenomenon known as the "neighbor explosion." This makes them unsuitable for large-scale graphs where the capture of extended neighborhoods is crucial. Layer-wise sampling, such as in FastGCN [CMX18], although designed to mitigate this explosion by sampling nodes per layer, often loses valuable connectivity information, leading to a sparse and inaccurate representation of the graph's multi-hop structure. Furthermore, subgraph sampling methods such as Cluster-GCN [CLS$^+$19] focus on localized computations and overlook vital inter-subgraph connections, failing to capture the broader topology necessary for a comprehensive multi-hop domain analysis.

GraphSnapShot represents a significant advancement in the field by addressing these inefficiencies. It introduces a hybrid sampling method that combines an initial comprehensive preprocessing of the graph with an intelligent dynamic sampling of multi-hop neighborhoods. This strategic approach allows for the maintenance of structural detail and computational efficiency, even as it scales to capture wider local topologies. GraphSnapShot dynamically adjusts its sampling based on the graph's evolving structure, which not only enhances the accuracy of the multi-hop neighborhood representation but also significantly reduces the computational load. Demonstrated by its considerable reduction in

computation time compared to traditional methods, GraphSnapShot provides a new strategy for efficient graph analysis in complex multi-hop domains.

4 Model Construction

Algorithm 1 Static Sampling and Dynamic ReSampling

1: **function** PREPROCESS(G)
2:     $sG \leftarrow$ STATICSAMPLE(G)
3:      **return** sG
4: **end function**
5: **function** DYNAMICSAMPLE(G, α, N)
6:     $num \leftarrow \alpha \times N$
7:     $sNodes \leftarrow$ SAMPLEFROMDISK(G, num)
8:      **return** $sNodes$
9: **end function**
10: **function** RESAMPLE(sG, α, N)
11:     $num \leftarrow \alpha \times N$
12:     $sNodes \leftarrow$ SAMPLEFROMMEMORY(sG, num)
13:      **return** $sNodes$
14: **end function**

Algorithm 2 Local Snapshot Swap
---
    function ADD(G, $Nodes$)
2:     $G \leftarrow G \cup Nodes$
    $Nodes \leftarrow$ MOVETOMEMORYFROMDISK($Nodes$)
4:      **return** G
    end function
6: **function** REMOVE(G, $Nodes$)
    $G \leftarrow G - Nodes$
8:     $Nodes \leftarrow$ MOVETODISKFROMMEMORY($Nodes$)
     return G
10: **end function**

Algorithm 3 GraphSnapShot Process
---
    function Process(G, α, N)

2:      $sG \leftarrow$ Preprocess(G)

      while ($InComputation$) **do**

4:         $dNodes \leftarrow$ DynamicSample(G, α, N)

         $rNodes \leftarrow$ ReSample(sG, $(1-\alpha)$, N)

6:         $uG \leftarrow$ Combine($dNodes$, $rNodes$)

         $sG \leftarrow$ Remove(sG, $rNodes$)

8:         $sG \leftarrow$ Add(sG, $dNodes$)

         $result \leftarrow$ Compute(uG)

10:      **end while**

      return $result$

12: **end function**

GraphSnapShot is designed to solve the inefficacy for current graph sampling methods on large graphs. It integrates the benefits of both static and dynamic processing. In the preprocessing phase, a static graph snapshot is sampled and stored in memory from the disk. Then, in the computation phase, the algorithm re-samples the graph snapshot, using a mix of data partially from memory and partially from disk. Concurrently, a smaller, dynamic snapshot is sampled from the disk. This dynamic snapshot is crucial as it is used to swap with the portion of the graph stored in memory. This process ensures the computation can always learn "fresh" snapshot from memory while not explode the memory. GraphSnapShot allows the algorithm to continuously update and adapt to changes in the graph while maintaining a manageable memory footprint, making it an effective tool for processing large, dynamic graph data.

Bridging the Gap Between Pure Dynamic Algorithms and Static Memory Storage with GraphSnapShot: Dynamic graph algorithms, such as GraphSAGE, which require resampling the entire graph for each computation, incur significant overhead. This is primarily due to the need for constant, full-graph resampling to accommodate the ever-changing nature of dynamic graphs. GraphSnapShot addresses this issue by establishing a static snapshot at preprocessing phase and use dynamic processing to update partial snapshot at computation. This approach significantly reduces the computational burden by providing a stable, memory-stored snapshot (sG) as a baseline. GraphSnapShot ensures efficient computation and maintains the flexibility needed to handle dynamic changes in the graph.

Tradeoff Between Dynamic Sampling and Resampling: In GraphSnapShot, the balance between dynamic sampling (*DynamicSample*) and resampling (*ReSample*) plays a pivotal role in maintaining an up-to-date representation of the graph. Dynamic sampling involves integrating new nodes from the disk into memory, capturing the latest updates and changes. Resampling, on the other hand, focuses on sampling and processing existing nodes from memory, which

are then rotated back to disk storage. However, if *DynamicSample* is too large, it is close to traditional sampling, resulting in prolonged disk retrieval times. Conversely, while resampling processes memory-stored nodes, an excessive *Re-Sample* may overly rely on memory data, potentially leading to a decrease in accuracy. Striking this balance is essential: moderating dynamic sampling to prevent inefficiencies and ensuring resampling maintains graph accuracy without an undue dependence on memory. In short, this is a tradeoff between optimizing memory use and accuracy.

Local Snapshot Swap Strategy: The *Add* and *Remove* functions are designed to simulate the swap of nodes between memory and disk, reflecting the dynamic shifts between active and less active parts of the graph. This local snapshot swapping strategy is crucial for handling large-scale dynamic graphs, as it reduces the demand on memory and computational resources.

In conclusion, the design of our GraphSnapShot Process algorithm considers the complexity and variability of dynamic graph analysis. By combining static sampling with dynamic resampling, the algorithm effectively manages and analyzes graphs whose structures evolve over time, demonstrating robust performance and flexibility in handling large-scale network analyses.
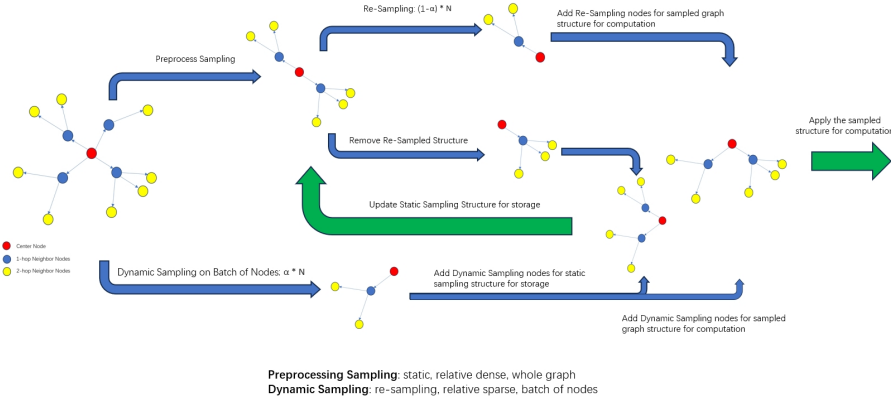

Figure 1: GraphSnapShot Model

5 Experimental Result

In our experiments, we explored the efficacy of the GraphSnapShot algorithm by testing the tradeoff between static and dynamic resampling. Static resampling (proportional to α) focuses on quick access to graph topology information stored in memory. Dynamic resampling (proportional to $1 - \alpha$), on the other hand, addresses the real-time changes in the graph, frequently requiring access to disk-stored data and swap in the memory to update the in-memory topology.

Figure 2: result of 1-hop expansion using the GraphSnapShot

Figure 3: result of 2-hop expansion using the GraphSnapShot

Figure 4: result of 3-hop expansion using the GraphSnapShot

Figure 5: result of 4-hop expansion using the GraphSnapShot

Analyzing the outcomes of these experiments provides insights into the algorithm's performance:

- **Accuracy and Alpha**: Our results showed a general increase in accuracy as alpha increased, which shows dynamic resampling has critical roles in enhancing model precision. A higher rate of dynamic resampling implies

a greater reliance on the stable structure of the graph, aiding in capturing patterns in graph topology. A relative high dynamic resampling rate is beneficial for the model's training stability and generalizability.

- **Loss and Alpha**: Loss decreased with an increase in alpha, further validating the importance of dynamic resampling in improving model performance. Lower loss indicates reduced error in data fitting, suggesting that static resampling helps the algorithm to learn and predict the graph's structure more accurately. This could also imply that dynamic resampling offers a stable learning environment, potentially reducing the risk of overfitting.

- **Alpha and Training Time**: Our experiments indicated the existence of an optimal alpha value between 0 and 1, which minimizes training time, and with a little loss in acuracy. We call this point as "tradeoff" point between static and dynamic resampling. Over-dependence on dynamic resampling (low alpha) could increase the computational burden due to frequent updates required for rapidly changing graph structures, while excessive reliance on static resampling (high alpha) might cause the model to overly learn from limited portion of the graph structure in memory, deduce its ability to effectively learn from the entire graph structure.

The effectiveness of the GraphSnapShot algorithm lies in how it balances static and dynamic resampling, which is controlled by α and $1 - \alpha$ respectively. This balance is key for handling the trade-off between using fast, in-memory data for stable and consistent performance (achieved through static resampling), and accessing data from disk to keep up with the latest updates (achieved through dynamic resampling). Fine-tuning the α parameter allows for calibration of the algorithm's focus between stability and adaptability. This adjustment is essential for achieving optimal performance across various graph analysis tasks. The GraphSnapShot excels in processing dynamic graph data and tasks that require exploration of graph structures, particularly for applications with complex and large-scale network structures.

References

[CLS+19] Wei-Lin Chiang, Xuanqing Liu, Si Si, Yang Li, Samy Bengio, and Cho-Jui Hsieh. Cluster-GCN. In *Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery &amp Data Mining.* ACM, jul 2019.

[CMX18] Jie Chen, Tengfei Ma, and Cao Xiao. Fastgcn: Fast learning with graph convolutional networks via importance sampling, 2018.

[HYL18] William L. Hamilton, Rex Ying, and Jure Leskovec. Inductive representation learning on large graphs, 2018.

A   Appendix

Dataset	Nodes	Edges	Features	Classes
PubMed	19,717	44,338	500	3
Cora	2,708	5,429	1,433	7
CiteSeer	3,312	4,732	3,703	6

Table 1: Comparison of PubMed, Cora, and CiteSeer in Terms of Nodes, Edges, Features, and Classes

Operation	Duration (seconds)	Simulation Frequency
Simulated Disk Read	5.0011	0.05
Simulated Disk Write	1.0045	0.05
Simulated Memory Access	0.0146	0.05

Table 2: Simulation Durations and Frequencies

Operation	k_hop_sampling	k_hop_retrieval	k_hop_resampling
Simulated Disk Read	✓		✓
Simulated Disk Write	✓		✓
Simulated Memory Access		✓	

Table 3: Function Access Patterns for Different Operations

Table 4: Experimental Settings - Setting 1

Dataset	Alpha	Presampled Nodes	Resampled Nodes	Sampled Depth
CiteSeer	0.1, 0.2, ..., 0.9	100	40	1, 2, 3, 4
Cora	0.1, 0.2, ..., 0.9	100	40	1, 2, 3, 4
PubMed	0.1, 0.2, ..., 0.9	100	40	1, 2, 3, 4

Table 5: Experimental Settings - Setting 2

Dataset	Alpha	Presampled Nodes	Resampled Nodes	Sampled Depth
CiteSeer	0.1, 0.2, ..., 0.9	20	10	1, 2, 3, 4
Cora	0.1, 0.2, ..., 0.9	20	10	1, 2, 3, 4
PubMed	0.1, 0.2, ..., 0.9	20	10	1, 2, 3, 4